S7 - Quiz

Due Jun 17 at 9am **Points** 100 **Questions** 8
Available Jun 10 at 1pm - Jun 17 at 9am **Time Limit** 30 Minutes

Instructions

Instructions:

1. You have 30 minutes to attempt the quiz
2. Once you start the quiz, you cannot go back and re-attempt it
3. You will not find answers online, so please make sure you are ready for the quiz
4. For Multiple Answer Questions, ALL the answers must be correct to score any point

Sometimes you might see multiple empty options. Please do not consider those empty options, that's some rendering issue, the options you see are the only options available for that question.

Attempt History

	Attempt	Time	Score
LATEST	Attempt 1	3 minutes	90 out of 100

Score for this quiz: **90** out of 100
Submitted Jun 11 at 3:45pm
This attempt took 3 minutes.

Question 1 — 10 / 10 pts

Select all which are true.

Dropout...

- [ ] Increases Test Acc while reducing Train Acc.
- [ ] Increases Train Accuracy
- [ ] Increases Text Accuracy

Correct!

- [x] Reduces the gap between TestAcc and TrainAcc

Question 2

10 / 10 pts

Batch Normalization should be added before the prediction layer.

- [ ] True

Correct!

- [x] False

Question 3

10 / 10 pts

Batch Normalization can be added only

Correct!

- [x] After at least one Convolution
- [ ] Even before the first Convolution

Question 4

10 / 10 pts

Select all which are true

- [ ] Adding LR Scheduler always increase accuracy

Correct!

- [x] Depending on GPU Bigger Batch size might speed up Epochs
- [ ] We do not need to maintain equal class representation in a batch

Correct!

- [x] To be on a safer side, it is always a good idea to shuffle the dataset.

Question 5

30 / 30 pts

Your model is overfitting. What all can be considered? (We'll retrain the model with the options that you'll select below)

- [ ] Going ahead with top_5 accuracies
- [x] Adding more training data (but not touching test images) — Correct!
- [x] Reducing number of kernels — Correct!
- [ ] Increasing number of kernels/parameters
- [ ] Increasing number of layers
- [x] Adding/Changing Image Augmentation strategies — Correct!
- [x] Adding Batch Normalization (if not added earlier) — Correct!
- [x] Adding DropOut — Correct!
- [ ] Changing Learning Rate
- [ ] Changing the Optimizer

Question 6

10 / 10 pts

The images in our dataset are of size 100x100. Currently, you are at a layer where the resolution is 7x7x512. Which the best option from below (as we covered in the 10 codes)?

- ( ) Dense Layer
- ( ) Larger Kernel Size to convert 7x7 to 1x1

Correct!

- Using GAP, followed by FC or 1x1 to match number of classes
- MaxPooling

Question 7

10 / 10 pts

The activations for classes A, B, and C before softmax were 10, 8, and 3.

The difference in softmax values for class A and class B would be

- 0.00085
- 12%

Correct!

- 76%
- 88%

Question 8

0 / 10 pts

The activations for classes A, B, and C before softmax were 10, 8, and 3.

There is only 1 image in the dataset and the class happens to be B. If we are using Negative Log-Likelihood Loss, the loss value right now is:

- -2.127731

ou Answered

- 0.127731

rrect Answer

- 2.127731

-0.127731

Quiz Score: **90** out of 100